



北京大學

Artificial Neural Networks: Advanced



人工智能引论 第11课


主讲人：刘家瑛

2019年4月15日

- *Slides modified from Feifei Li.*



- Artificial Neurons
 - **Perceptrons** can do simple classification tasks
 - **Multi-Layer Non-Linear** Neural Network can approximate a wide range of functions
 - **Back-Propagation** with the chain rule to train the network
- Convolutional Neural Networks
 - Convolution operates on **local spatial** area
 - **Fewer parameters** compared with fully connected networks
 - **Automatically** extracts **features** from images



- CNN Architectures
 - AlexNet
 - VGG
 - GoogLeNet
 - ResNet
- Recurrent Neural Network
 - Vanilla RNN
 - Backpropagation through time
 - Long Short-Term Memory
- Beyond CNN and RNN
 - Unsupervised Learning
 - Generative Adversarial Network




- **CNN Architectures**
 - **AlexNet**
 - **VGG**
 - **GoogLeNet**
 - **ResNet**
- Recurrent Neural Network
 - Vanilla RNN
 - Backpropagation through time
 - Long Short-Term Memory
- Beyond CNN and RNN
 - Unsupervised Learning
 - Generative Adversarial Network



Architecture:

CONV1 MAX POOL1 NORM1 CONV2 MAX POOL2 NORM2 CONV3 CONV4
CONV5 MAX POOL3 FC6 FC7 FC8



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)

Details / Retrospectives:

- First use of ReLU
- Used Norm layers (not common anymore)
- Heavy data augmentation
- Dropout 0.5
- Batch size 128
- SGD Momentum 0.9
- Learning rate 1e-2, reduced by 10 manually when val accuracy plateaus
- L2 weight decay 5e-4
- 7 CNN ensemble: 18.2% → 15.4%

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons


[4096] **FC7**: 4096 neurons


[1000] **FC8**: 1000 neurons (class scores)

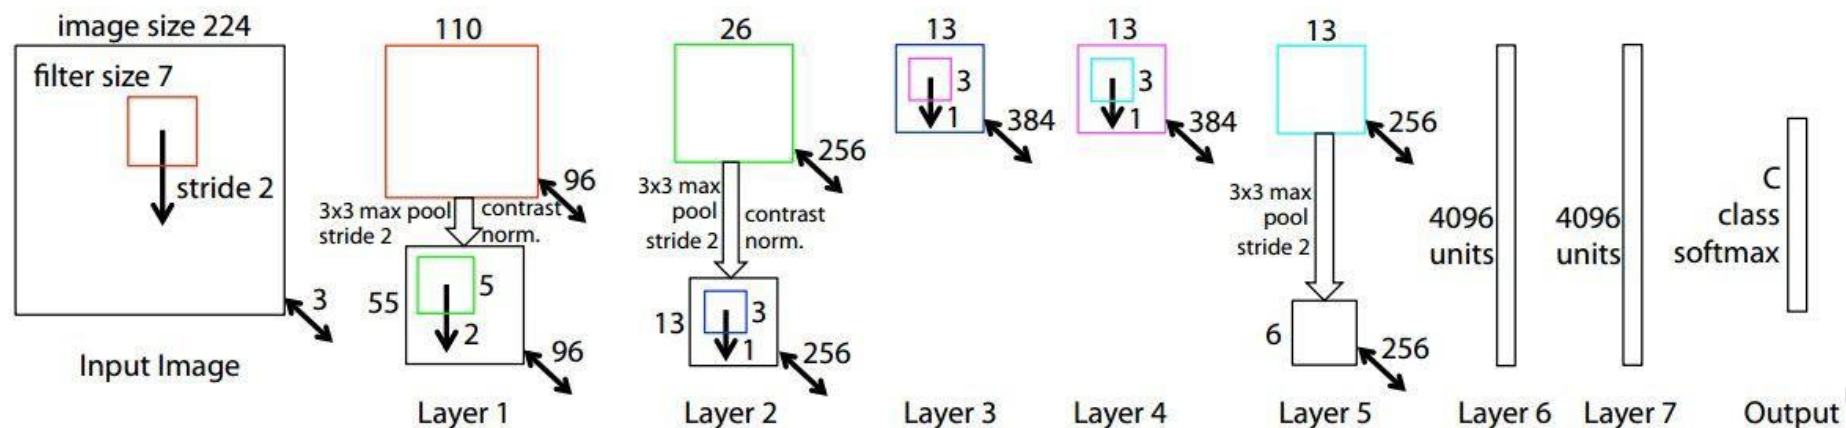
Details/Retrospectives:

- First use of ReLU
- Used Norm layers (not common anymore)
- Heavy data augmentation
- Dropout 0.5
- Batch size 128
- SGD Momentum 0.9
- Learning rate 1e-2, reduced by 10 manually when val accuracy plateaus
- L2 weight decay 5e-4
- 7 CNN ensemble: 18.2% → 15.4%

ILSVRC Winners








AlexNet but:

CONV1: change from (11x11 stride 4) x48 to (7x7 stride 2) x96

CONV3,4,5: instead of 384, 384, 256 filters use 512, 1024, 512

ImageNet top 5 error: 16.4% → 11.7%

ILSVRC Winners



Small filters, deeper networks

8 layers (AlexNet)


→ 16 - 19 layers (VGG16Net)

Only 3x3 CONV: stride 1, pad 1 and 2x2


MAX POOL: stride 2

11.7% top 5 error in ILSVRC'13 (ZFNet)

→ 7.3% top 5 error in ILSVRC'14



AlexNet



VGG16


VGG19

Q: Why use smaller filters? (3x3 conv)

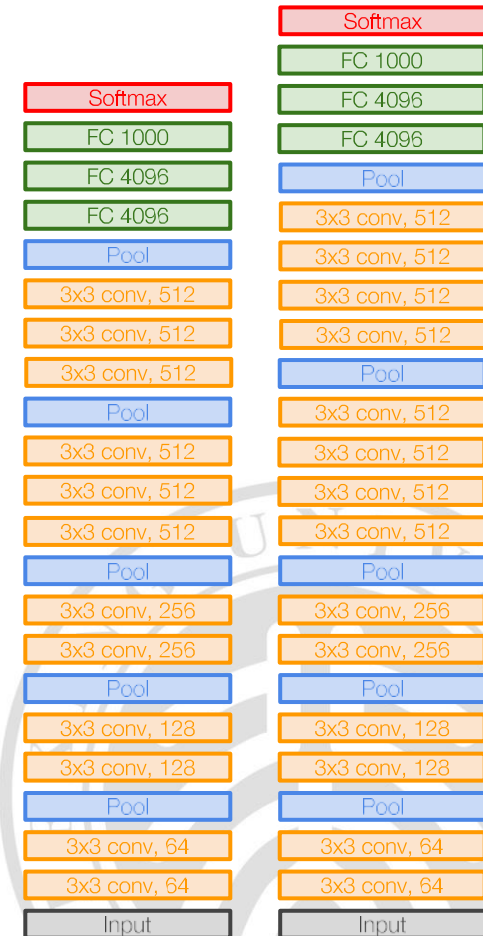
Stack of three 3x3 conv (stride 1) layers has same **effective receptive field** as one 7x7 conv layer

But deeper, more non-linearities

And fewer parameters: $3 * (3^2 C^2)$ vs. $7^2 C^2$ for C channels per layer



AlexNet



VGG16

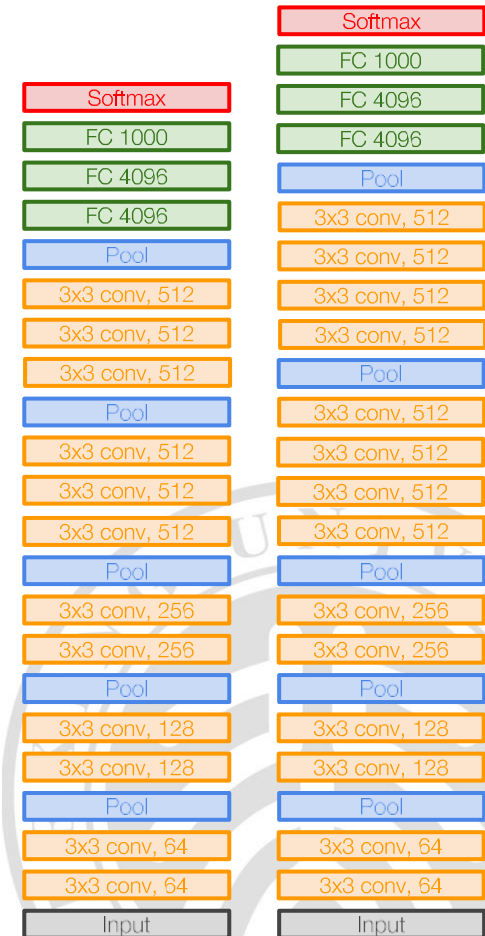
VGG19

- Details:**

- ILSVRC'14 2nd in classification, 1st in localization
- Similar training procedure as Krizhevsky 2012
- No Local Response Normalisation (LRN)
- Use VGG16 or VGG19 (VGG19 only slightly better, more memory)
- Use ensembles for best results
- FC7 features generalize well to other tasks




AlexNet



VGG16


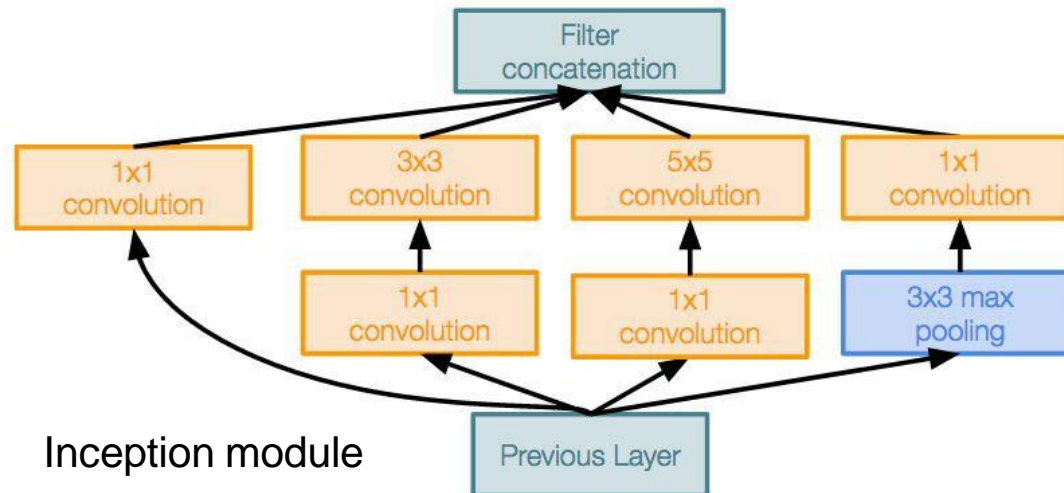
VGG19

ILSVRC Winners





Deeper networks, with computational efficiency


- 22 layers
- Efficient “Inception” module
- No FC layers
- Only 5 million parameters!
12x less than AlexNet
- ILSVRC’14 classification winner (6.7% top 5 error)



“Inception module”:

Design a good local network topology (network within a network),
and then stack these modules on top of each other.





Naive Inception module

Apply parallel filter operations on the input from previous layer:

- Multiple receptive field sizes for convolution (1x1, 3x3, 5x5)
- Pooling operation (3x3)

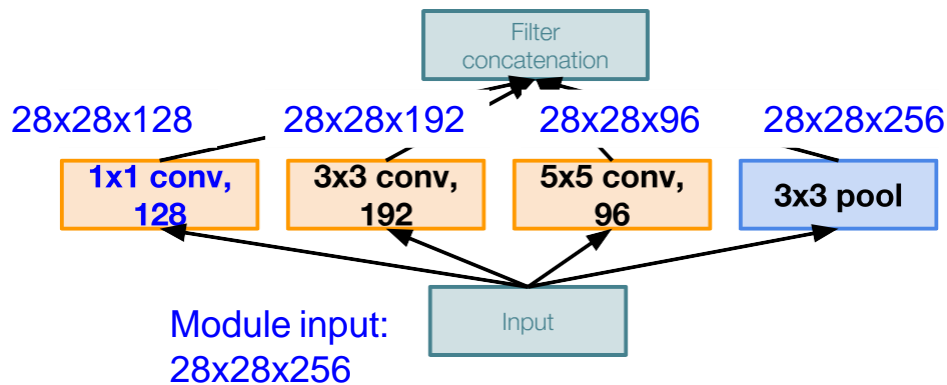
Concatenate all filter outputs together depth-wise

Q: What is the problem with this?
[Hint: Computational complexity]

Example:

Q3: What is output size after filter concatenation?

$$28 \times 28 \times (128 + 192 + 96 + 256) = 28 \times 28 \times 672$$



Naive Inception module

Q: What is the problem with this?
[Hint: Computational complexity]

Conv Ops:

[1x1 conv, 128] $28 \times 28 \times 128 \times 1 \times 1 \times 256$

[3x3 conv, 192] $28 \times 28 \times 192 \times 3 \times 3 \times 256$

[5x5 conv, 96] $28 \times 28 \times 96 \times 5 \times 5 \times 256$

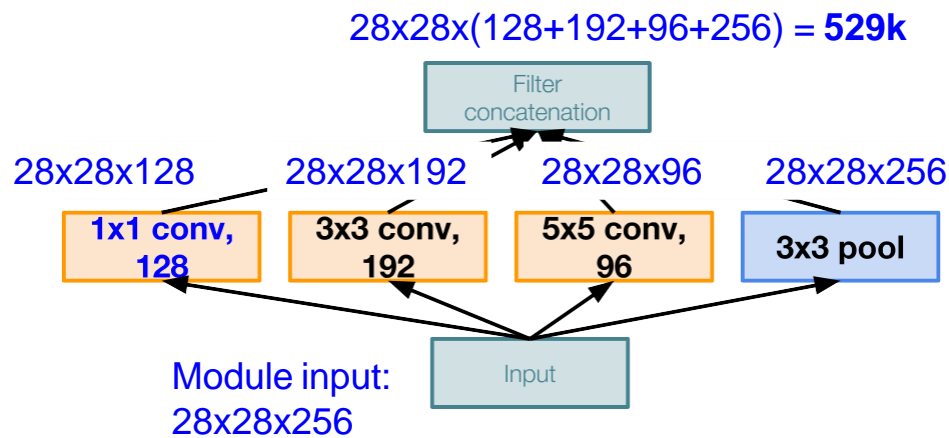
Total: 854M ops

Very expensive compute

Pooling layer also preserves feature depth, which means total depth after concatenation can only grow at every layer!

Example:

Q3: What is output size after filter concatenation?





Naive Inception module

Q: What is the problem with this?
[Hint: Computational complexity]


Solution: “bottleneck” layers that use 1×1 convolutions to reduce feature depth




Reminder: 1x1 Convolution



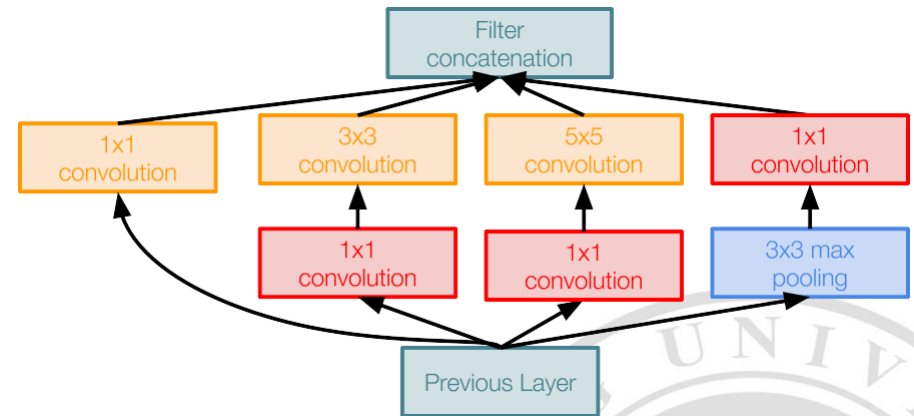
Reminder: 1x1 Convolution



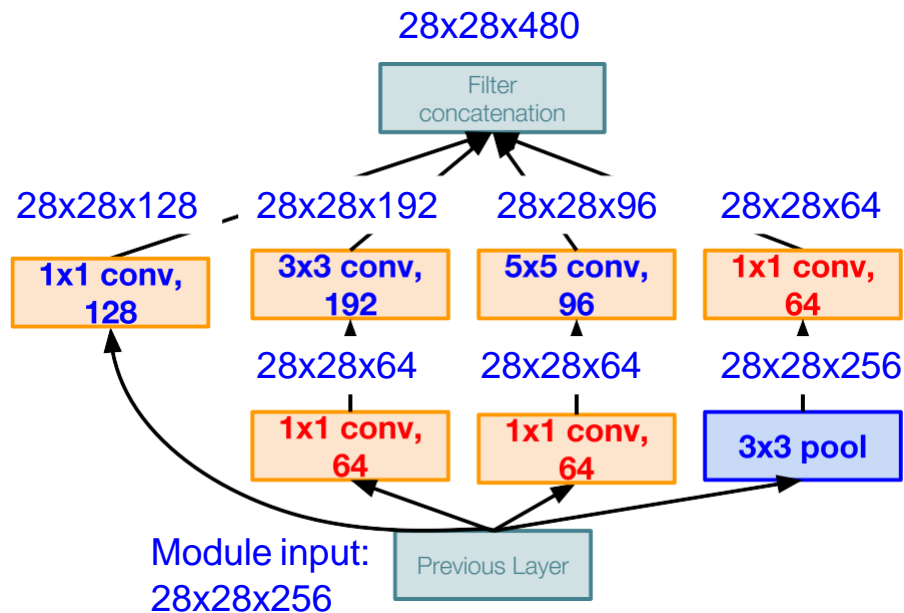
1x1 conv “bottleneck” layers



Naive Inception module



Inception module with dimension reduction



Inception module with dimension reduction

Using same parallel layers as naive example, and adding “1x1 conv, 64 filter” bottlenecks:

Conv Ops:

[1x1 conv, 64] 28x28x64x1x1x256

[1x1 conv, 64] 28x28x64x1x1x256

[1x1 conv, 128] 28x28x128x1x1x256

[3x3 conv, 192] 28x28x192x3x3x64


[5x5 conv, 96] 28x28x96x5x5x64

[1x1 conv, 64] 28x28x64x1x1x256

Total: 358M ops


Compared to 854M ops for naive version, bottleneck can also reduce depth after pooling layer

Full GoogLeNet architecture




Stem Network:
Conv-Pool-
2x Conv-Pool

Full GoogLeNet architecture




Stacked Inception Modules

Full GoogLeNet architecture




Full GoogLeNet architecture




Auxiliary classification outputs to inject additional gradient at lower layers
(AvgPool-1x1Conv-FC-FC-Softmax)

Full GoogLeNet architecture





22 total layers with weights

(parallel layers count as 1 layer => 2 layers per Inception module. Don't count auxiliary output layers)




Very deep networks using residual connections

- 152-layer model for ImageNet
- ILSVRC'15 classification winner (3.57% top 5 error)
- Swept all classification and detection competitions in ILSVRC'15 and COCO'15!



What happens when we continue stacking deeper layers on a “plain” convolutional neural network?



56-layer model performs worse on both training and test error
→ The deeper model performs worse, but it's not caused by overfitting!



Hypothesis: the problem is an *optimization* problem, deeper models are harder to optimize

The deeper model should be able to perform at least as well as the shallower model.


A solution by construction is copying the learned layers from the shallower model and setting additional layers to identity mapping.



Solution: Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping



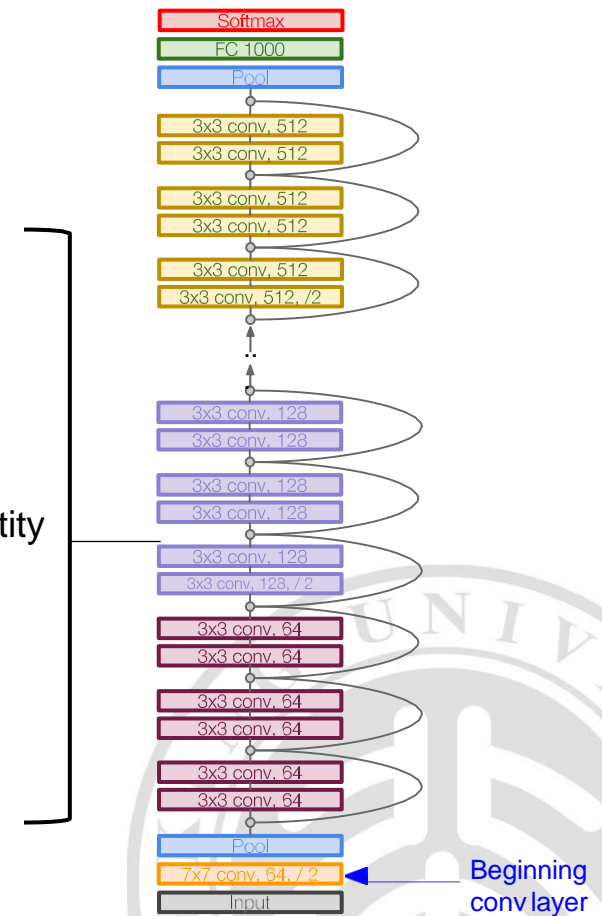

Solution: Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping



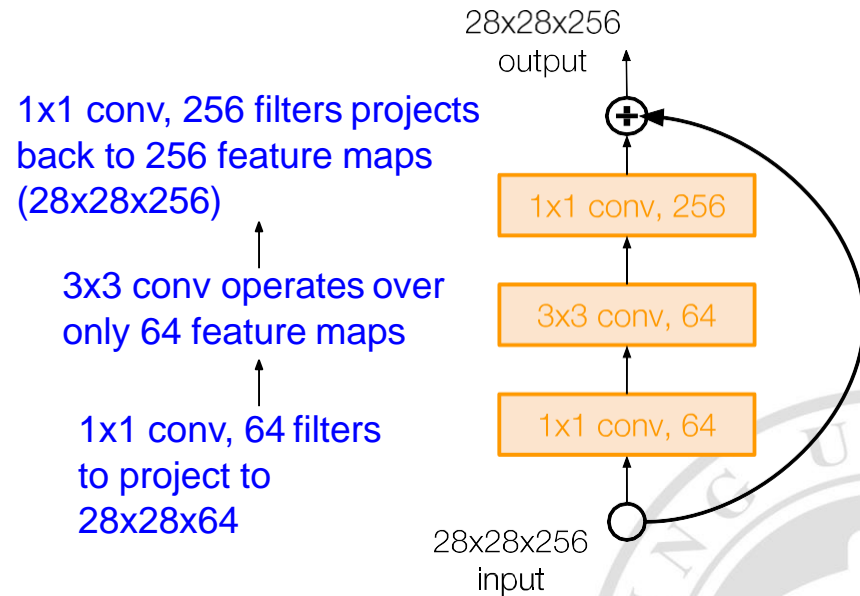
Use layers to fit residual
 $F(x) = H(x) - x$
 instead of
 $H(x)$ directly

Full ResNet architecture:

- Stack residual blocks
- Every residual block has two 3x3 conv layers
- Periodically, double # of filters and downsample spatially using stride 2 (/2 in each dimension)
- Additional conv layer at the beginning



For deeper networks
(ResNet-50+),
use “bottleneck” layer to
improve efficiency
(similar to GoogLeNet)



Training ResNet in practice:

- Batch Normalization after every CONV layer
- Xavier 2/ initialization from He et al.
- SGD + Momentum (0.9)
- Learning rate: 0.1, divided by 10 when validation error plateaus
- Mini-batch size 256
- Weight decay of $1e-5$
- No dropout used



Experimental Results

- Able to train very deep networks without degrading (152 layers on ImageNet, 1202 on Cifar)
- Deeper networks now achieve lower training error as expected
- Swept 1st place in all ILSVRC and COCO 2015 competitions

MSRA @ ILSVRC & COCO 2015 Competitions

- **1st places in all five main tracks**

- ImageNet Classification: “*Ultra-deep*” (quote Yann) **152-layer** nets
- ImageNet Detection: **16%** better than 2nd
- ImageNet Localization: **27%** better than 2nd
- COCO Detection: **11%** better than 2nd
- COCO Segmentation: **12%** better than 2nd

ILSVRC 2015 classification winner (3.6% top 5 error)
-- better than “human performance”!
(Russakovsky 2014)


- **VGG, GoogLeNet, ResNet** all in wide use, available in model zoos
- ResNet current best default
- Trend towards extremely deep networks
- Significant research centers around design of layer / skip connections and improving gradient flow
- Efforts to investigate necessity of depth Vs. width and residual connections




- CNN Architectures
 - AlexNet
 - VGG
 - GoogLeNet
 - ResNet
- **Recurrent Neural Network**
 - **Vanilla RNN**
 - **Backpropagation through time**
 - **Long Short-Term Memory**
- Beyond CNN and RNN
 - Unsupervised Learning
 - Generative Adversarial Network




one to one




one to many




many to one




many to many




many to many





 e.g. **Image Captioning**
 image → sequence of words




one to one




one to many




many to one



many to many




many to many




e.g. **Sentiment Classification**
sequence of words → sentiment


one to one




one to many




many to one




many to many




many to many





 e.g. **Machine Translation**
 seq of words → seq of words


one to one




one to many




many to one




many to many




many to many



e.g. Video classification on frame level





We can process a sequence of vectors \mathbf{x} by applying a **recurrence formula** at every time step:


$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state

some function with parameters W

old state


input vector at some time step




We can process a sequence of vectors \mathbf{x} by applying a **recurrence formula** at every time step:

$$h_t = f_W(h_{t-1}, x_t)$$

Notice: the same function and the same set of parameters are used at every time step.



The state consists of a single “*hidden*” vector \mathbf{h} :




$$h_t = f_W(h_{t-1}, x_t)$$




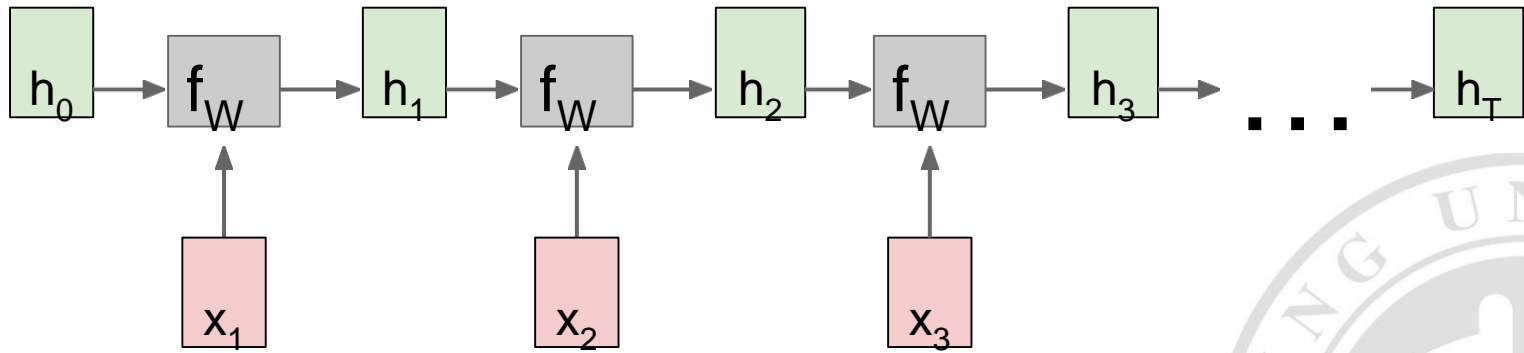
$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

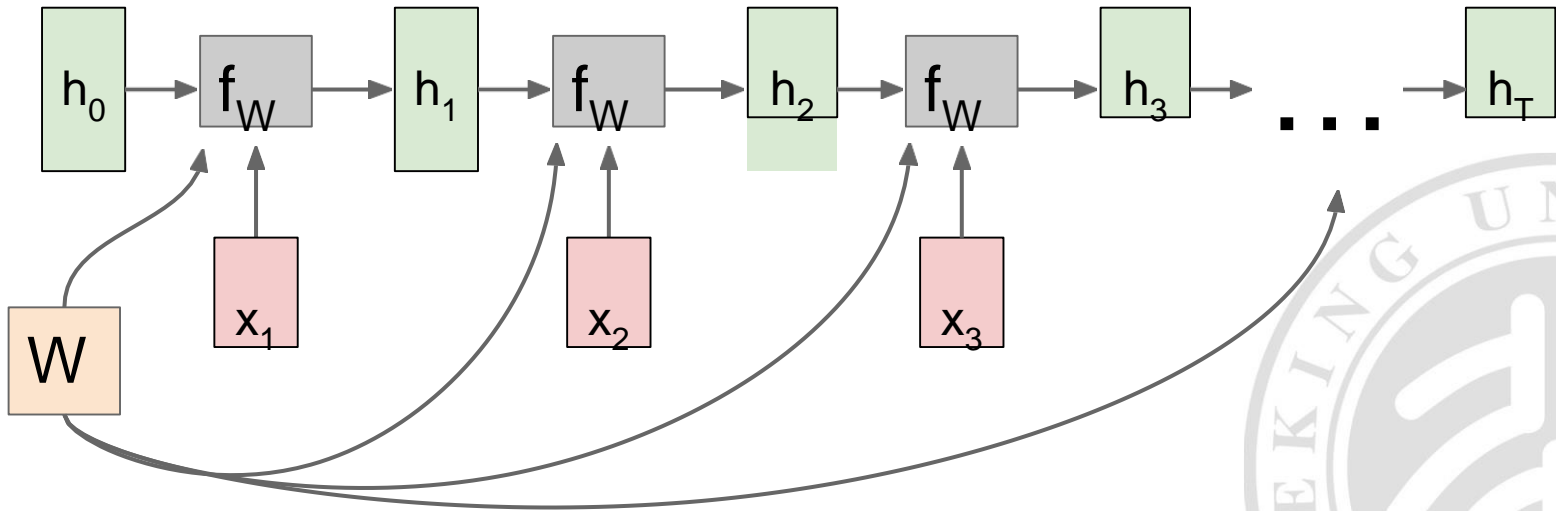
Sometimes called a “Vanilla RNN” or an “Elman RNN” after Prof. Jeffrey Elman




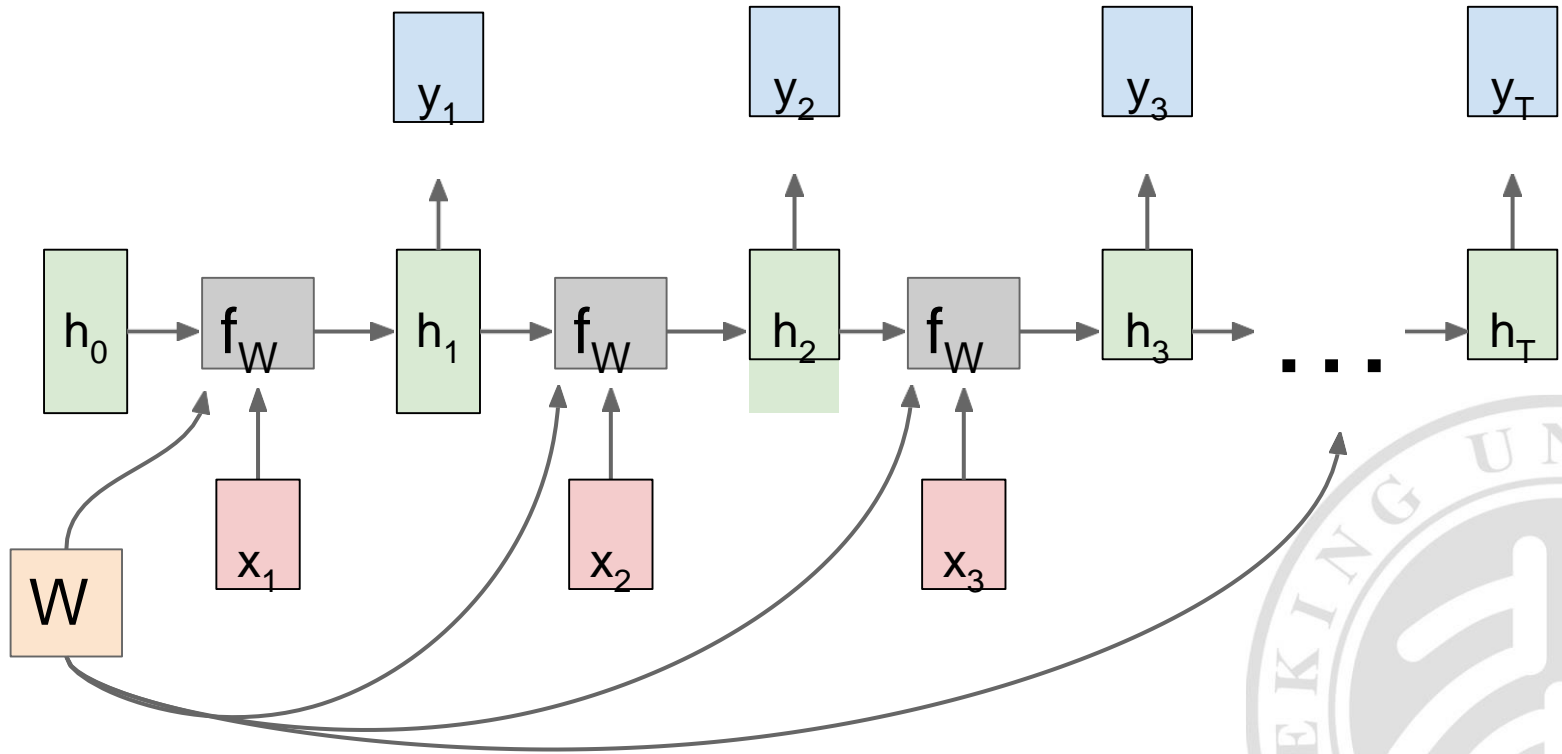





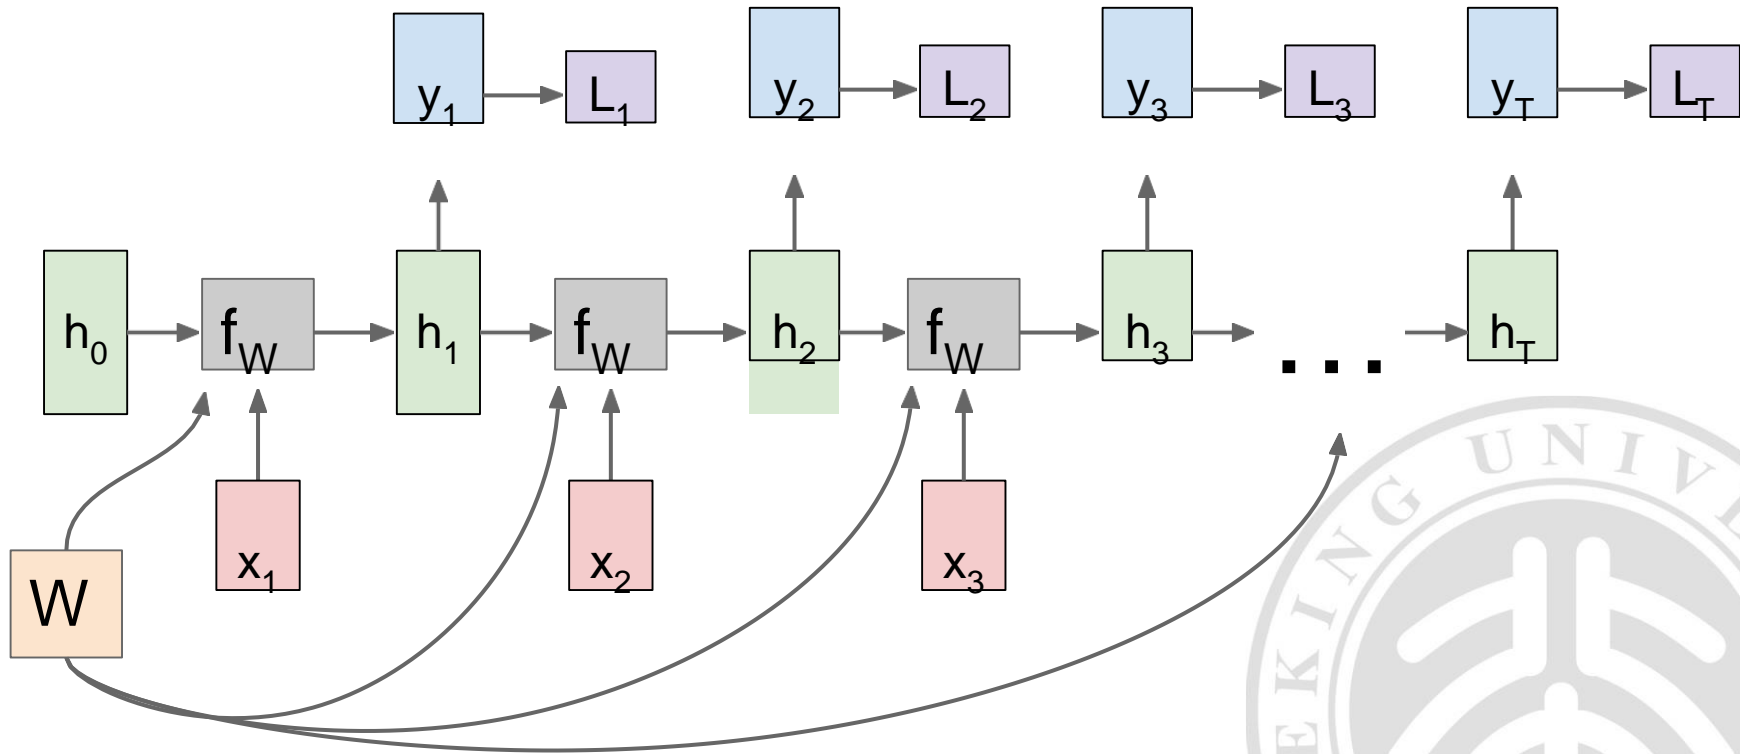
Re-use the same weight matrix at every time-step



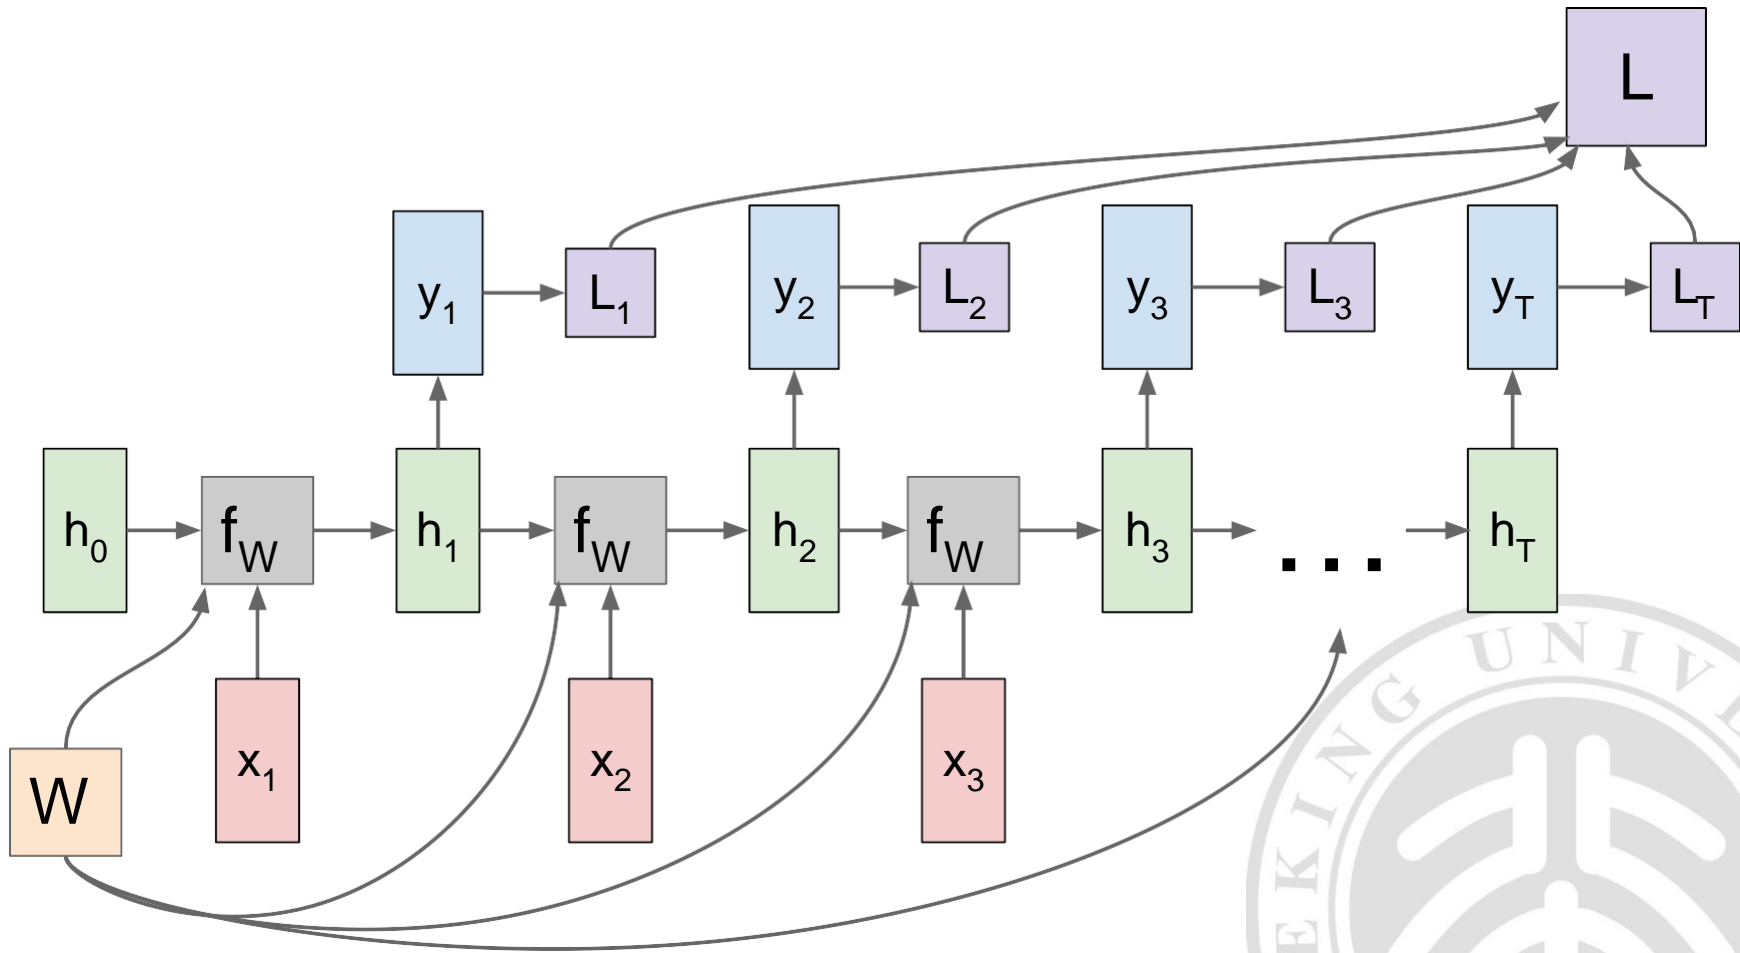
Many to Many




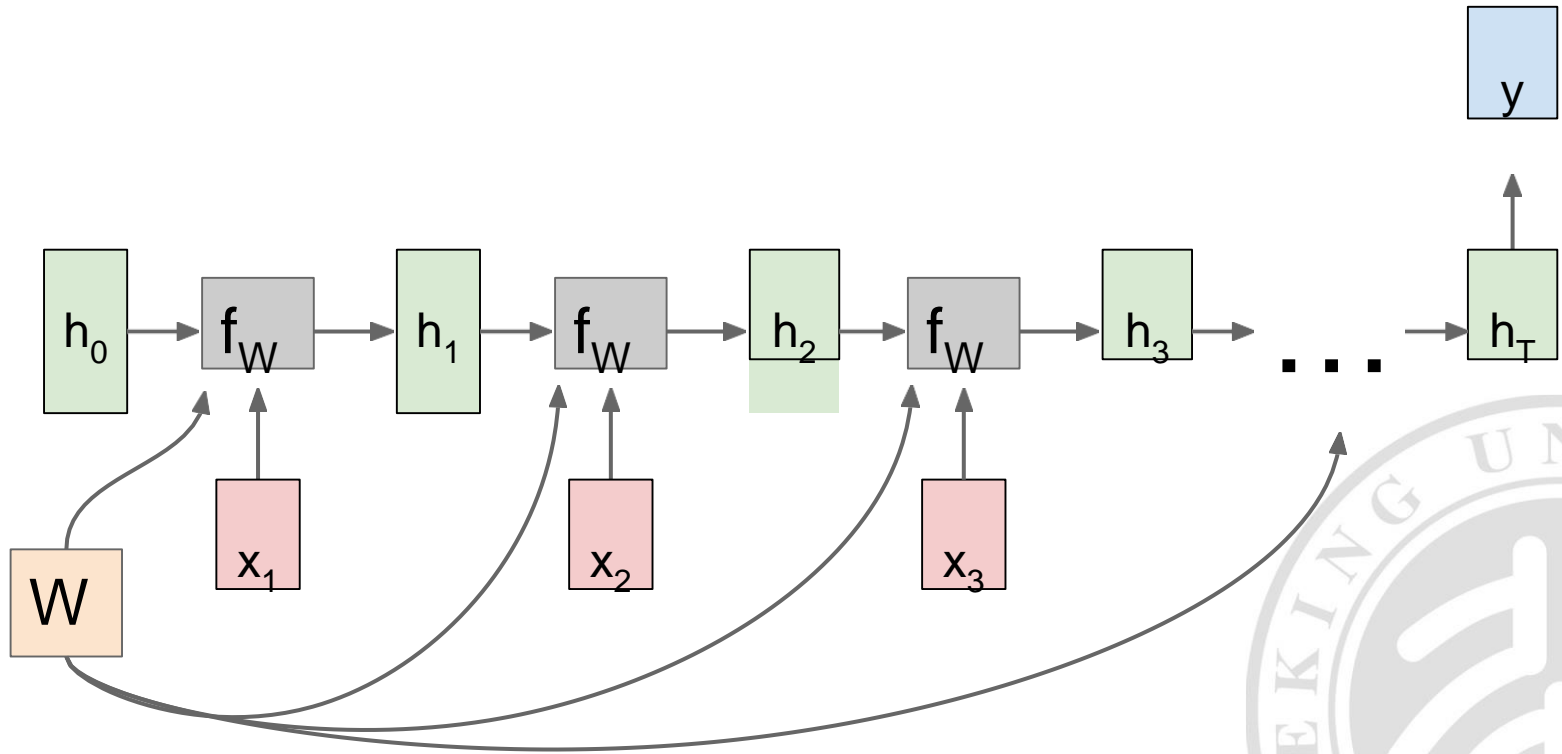
Many to Many




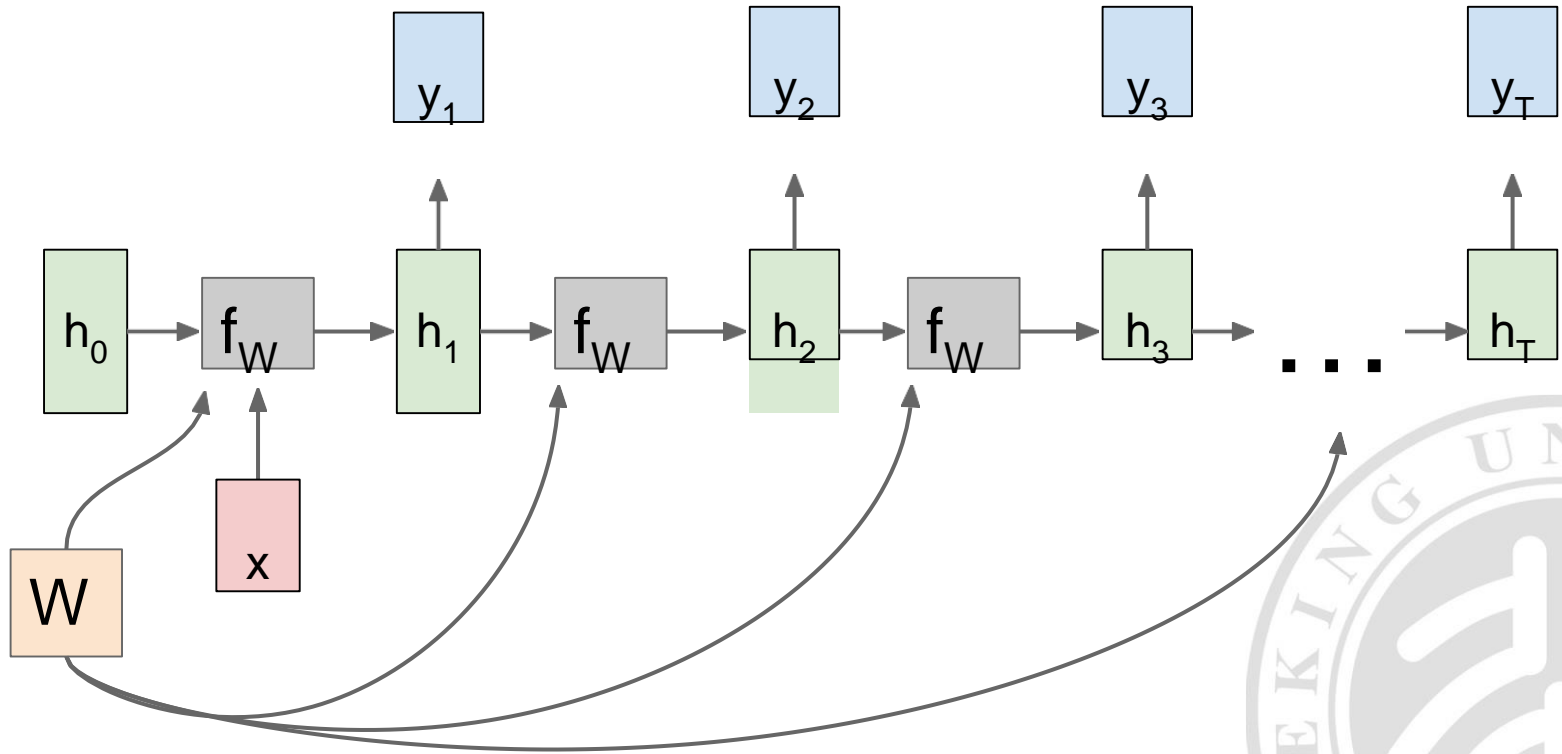
Many to Many




Many to One



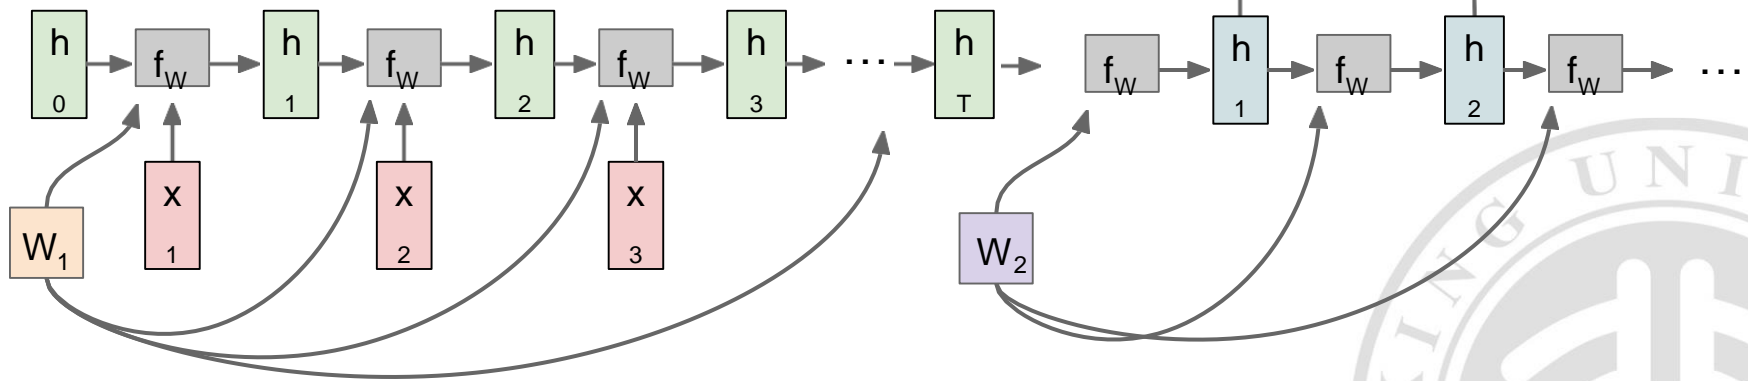
One to Many




Many to one: Encode input sequence in a single vector



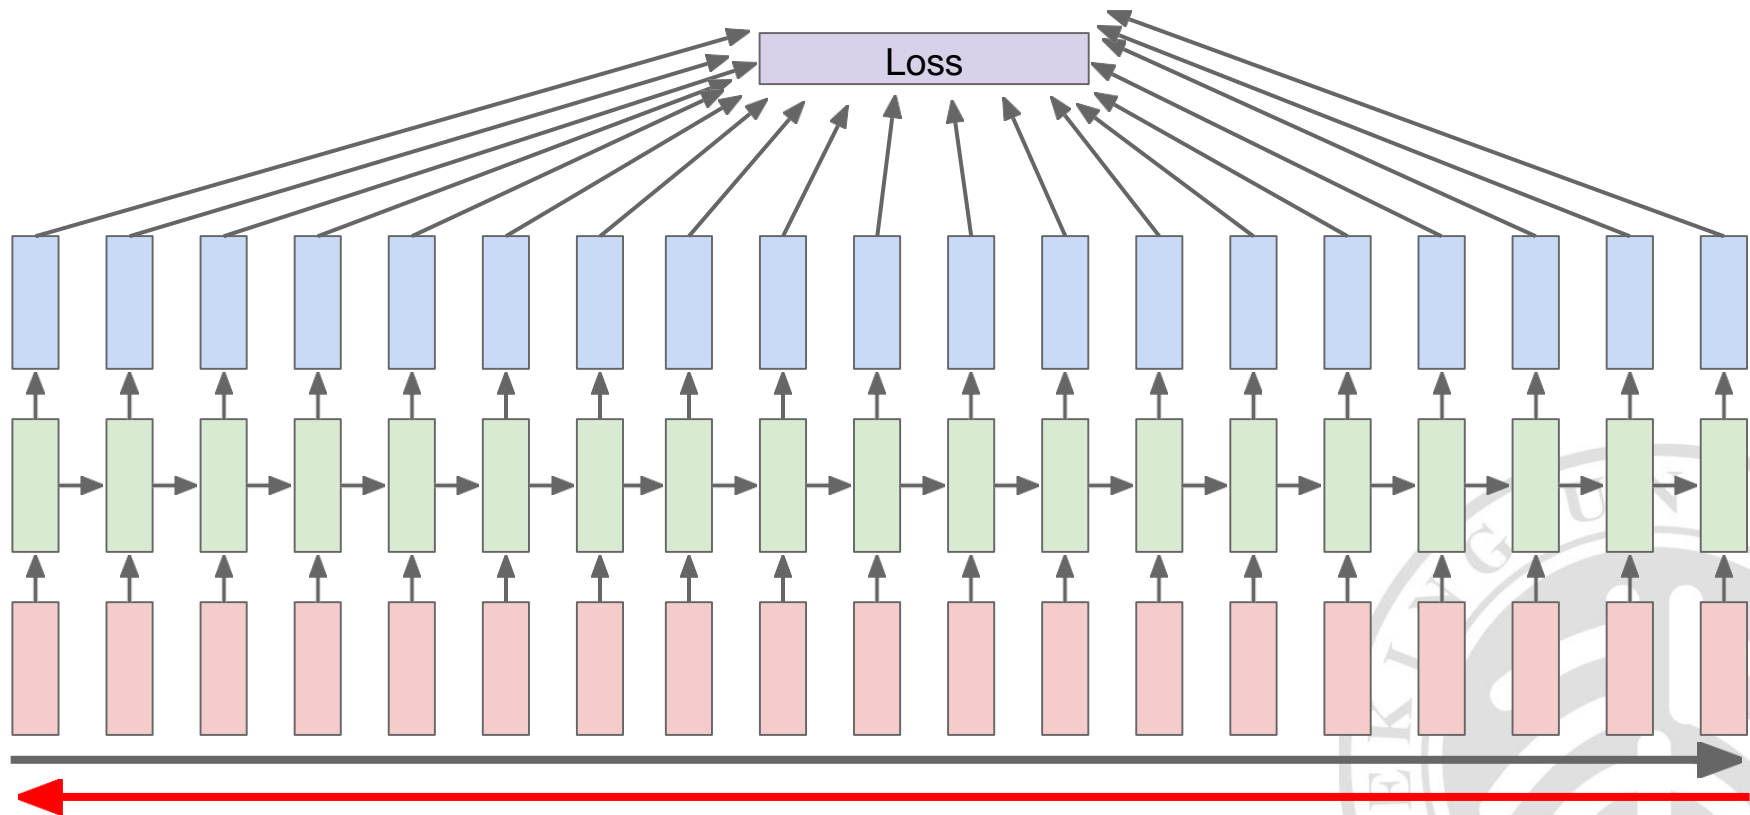
Many to one: Encode input sequence in a single vector



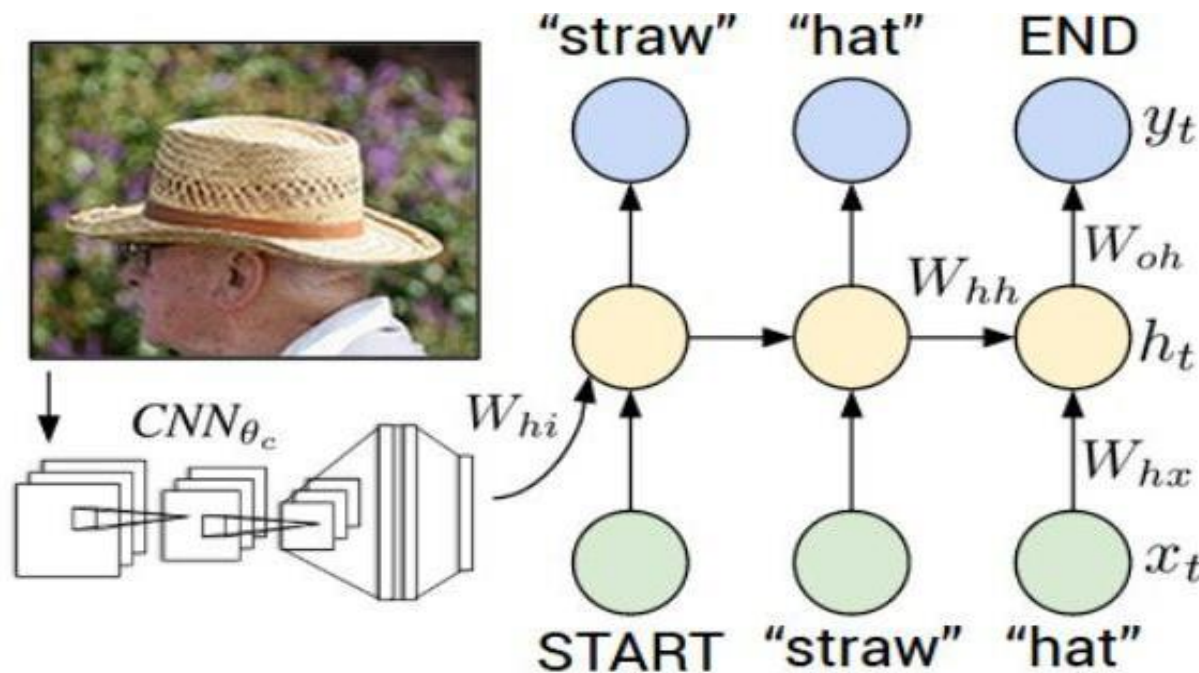
One to many: Produce output sequence from single input vector



Forward through entire sequence to compute loss, then backward through entire sequence to compute gradient



Example: Image Captioning



Explain Images with Multimodal Recurrent Neural Networks, Mao et al.

Deep Visual-Semantic Alignments for Generating Image Descriptions, Karpathy and Fei-Fei

Show and Tell: A Neural Image Caption Generator, Vinyals et al.

Long-term Recurrent Convolutional Networks for Visual Recognition and Description, Donahue et al.

Learning a Recurrent Visual Representation for Image Caption Generation, Chen and Zitnick

image

conv-64

conv-64

maxpool

conv-128

conv-128

maxpool

conv-256

conv-256

maxpool

conv-512

conv-512

maxpool

conv-512

conv-512

maxpool

FC-4096

FC-4096

V

Wih

y0

h0

x0
<STA
RT>

<START>



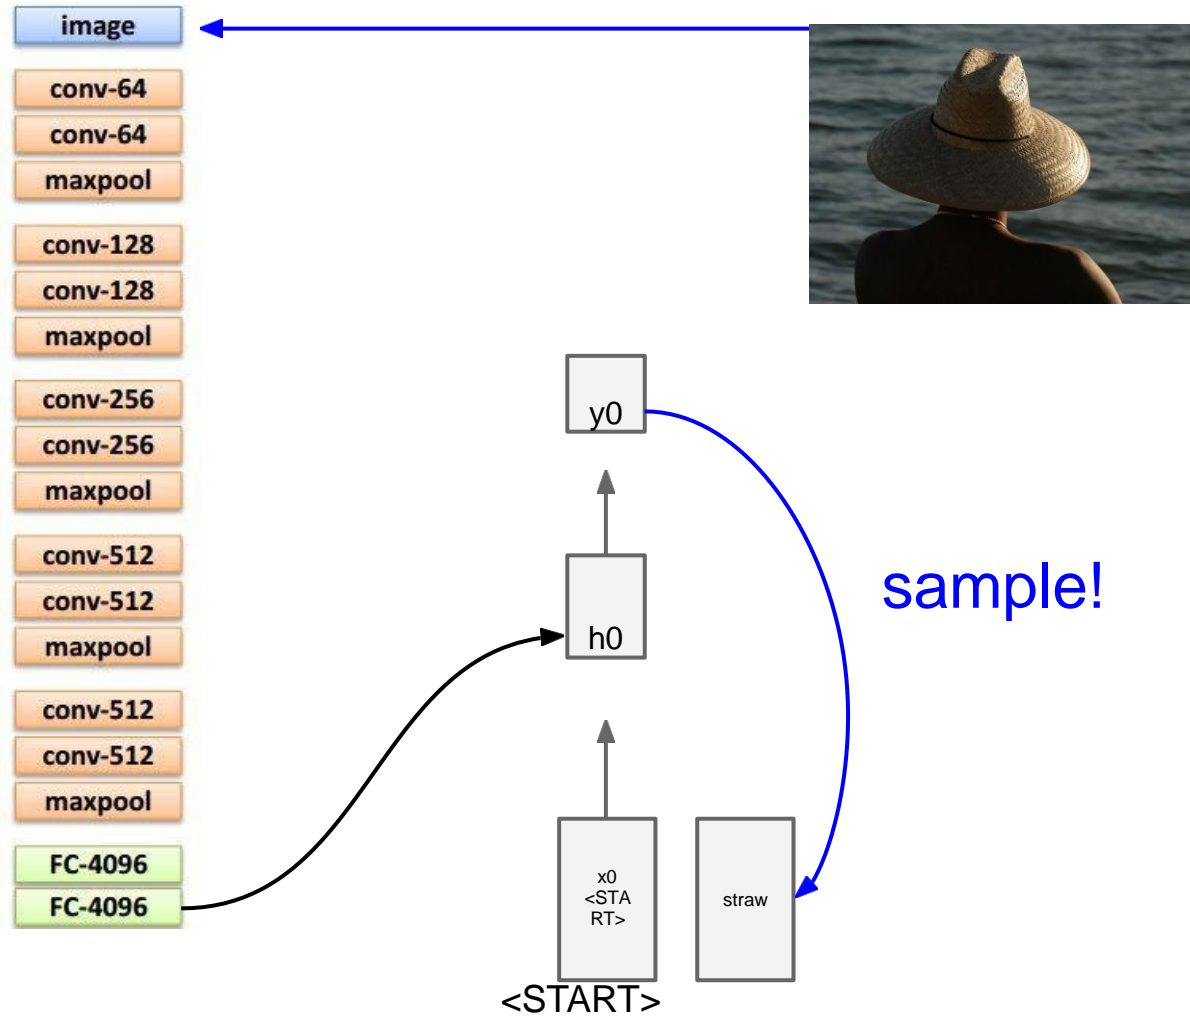
test image

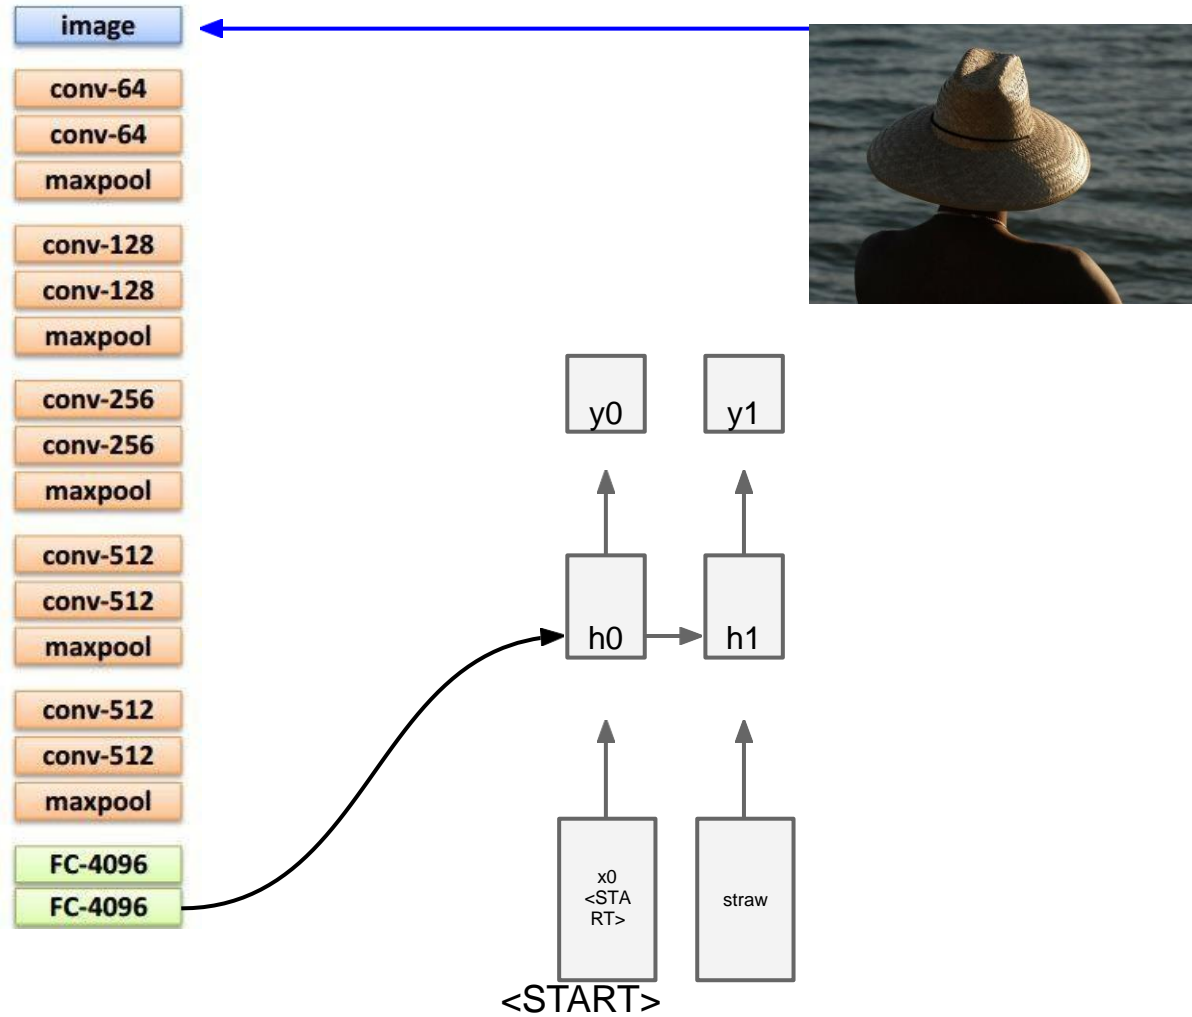
before:

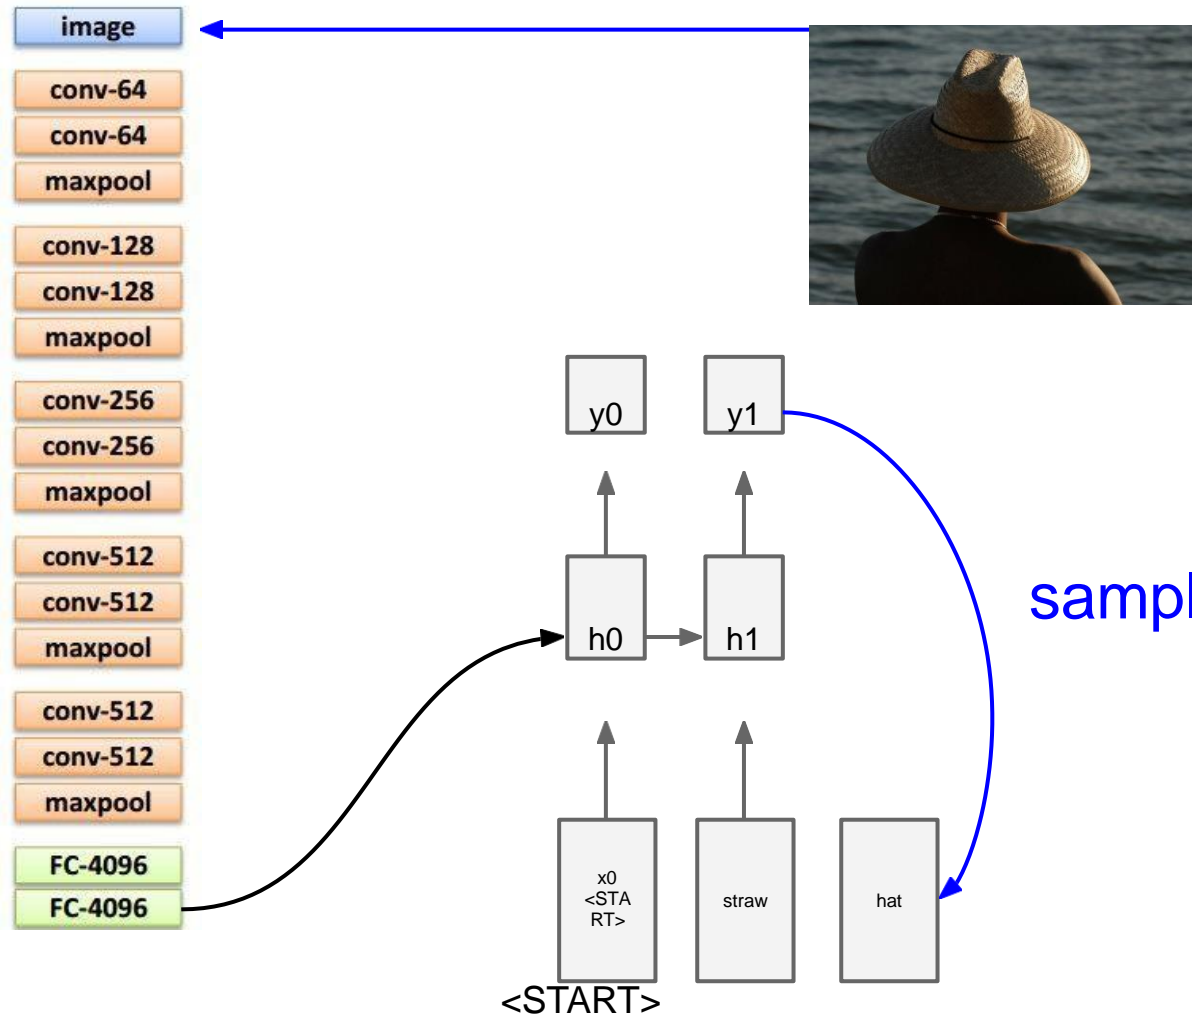
$$h = \tanh(W_{xh} * x + W_{hh} * h)$$

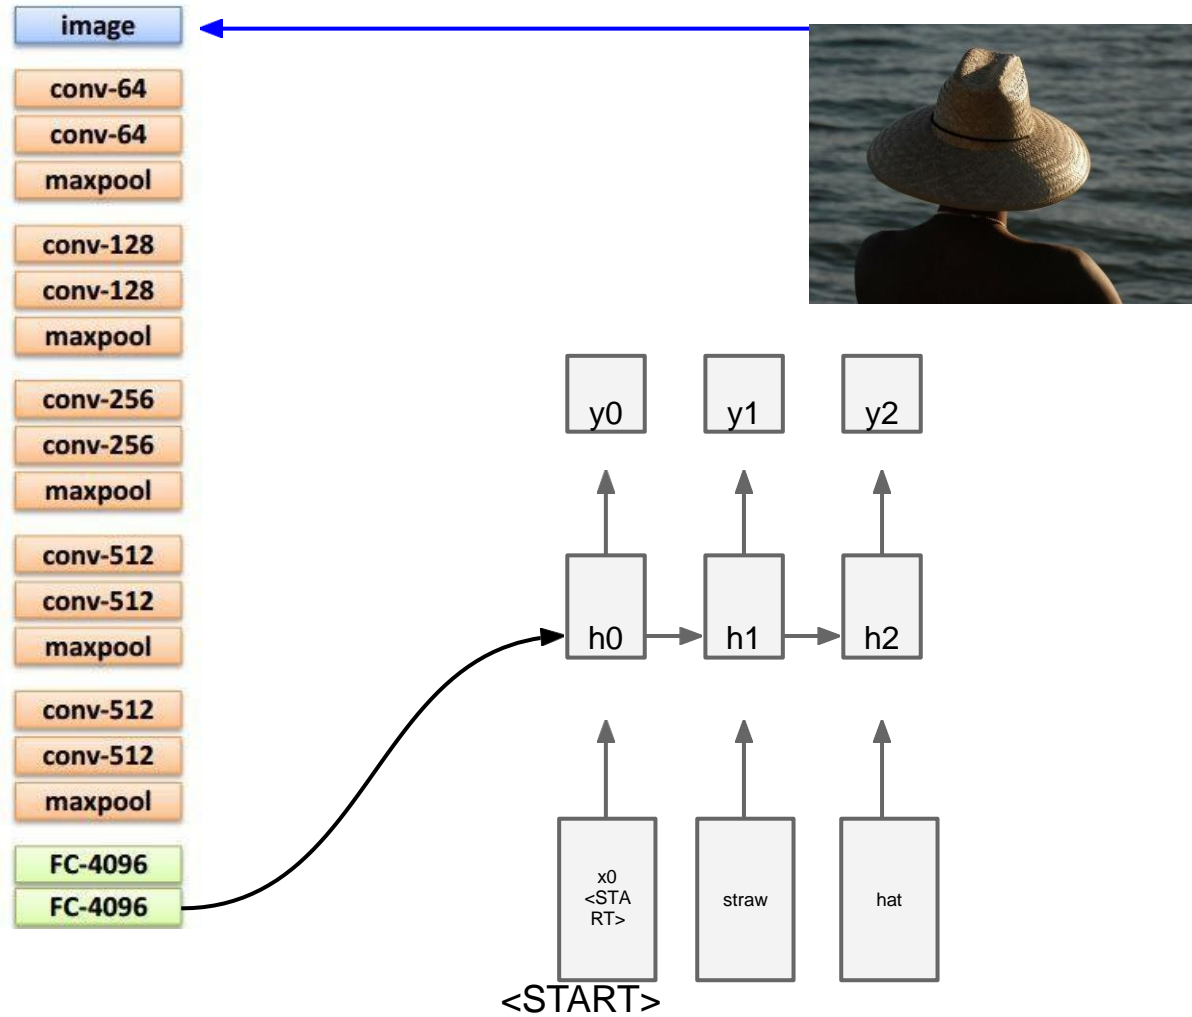
now:

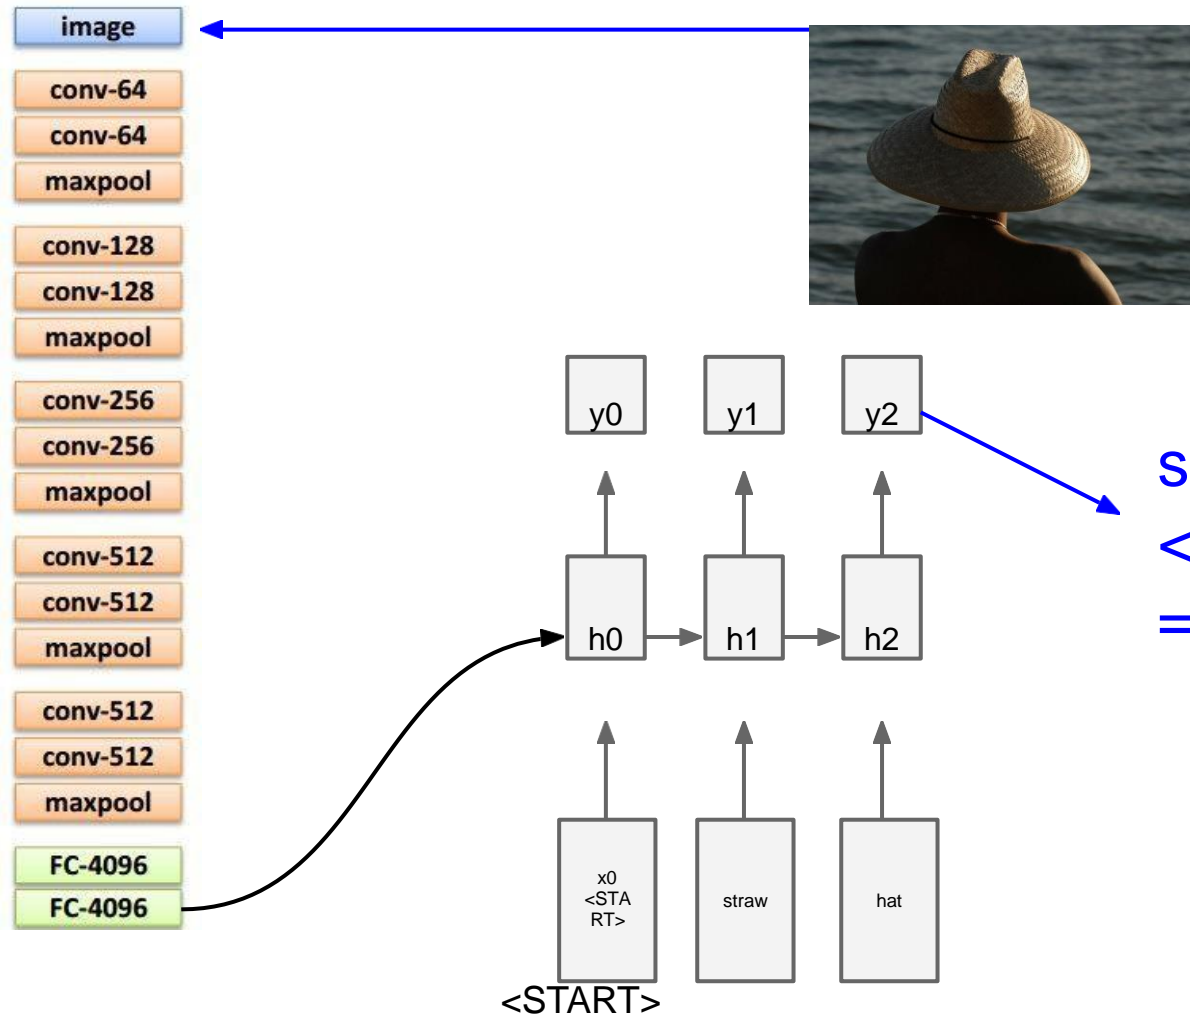
$$h = \tanh(W_{xh} * x + W_{hh} * h + W_{ih} * v)$$













A cat sitting on a suitcase on the floor



A cat is sitting on a tree branch



A dog is running in the grass with a frisbee



A white teddy bear sitting in the grass



Two people walking on the beach with surfboards



A tennis player in action on the court



Two giraffes standing in a grassy field



A man riding a dirt bike on a dirt track



A woman is holding a cat in her hand



A woman standing on a beach holding a surfboard



A bird is perched on a tree branch



A person holding a computer mouse on a desk



A man in a baseball uniform throwing a ball

$$h_t^l = \tanh W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$

$h \in \mathbb{R}^n$ $W^l [n \times 2n]$


LSTM:

$W^l [4n \times 2n]$


$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \text{sigm} \\ \text{sigm} \\ \text{sigm} \\ \text{tanh} \end{pmatrix} W^l \begin{pmatrix} h_t^{l-1} \\ h_{t-1}^l \end{pmatrix}$$

$$c_t^l = f \odot c_{t-1}^l + i \odot g$$

$$h_t^l = o \odot \tanh(c_t^l)$$

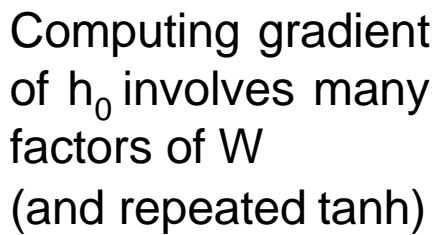


Backpropagation from h_t to h_{t-1}
multiplies by W
(actually W_{hh}^T)



$$\begin{aligned} h_t &= \tanh(W_{hh}h_{t-1} + W_{hx}x_t) \\ &= \tanh\left((W_{hh} \quad W_{hx}) \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}\right) \\ &= \tanh\left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}\right) \end{aligned}$$

Bengio et al, "Learning long-term dependencies with gradient descent is difficult", IEEE Transactions on Neural Networks, 1994
Pascanu et al, "On the difficulty of training recurrent neural networks", ICML 2013



Largest singular value < 1 :
Vanishing gradients

```
grad_norm = np.sum(grad * grad)
if grad_norm > threshold:
    grad *= (threshold / grad_norm)
```

74

Vanilla RNN

$$h_t = \tanh \left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix} \right)$$

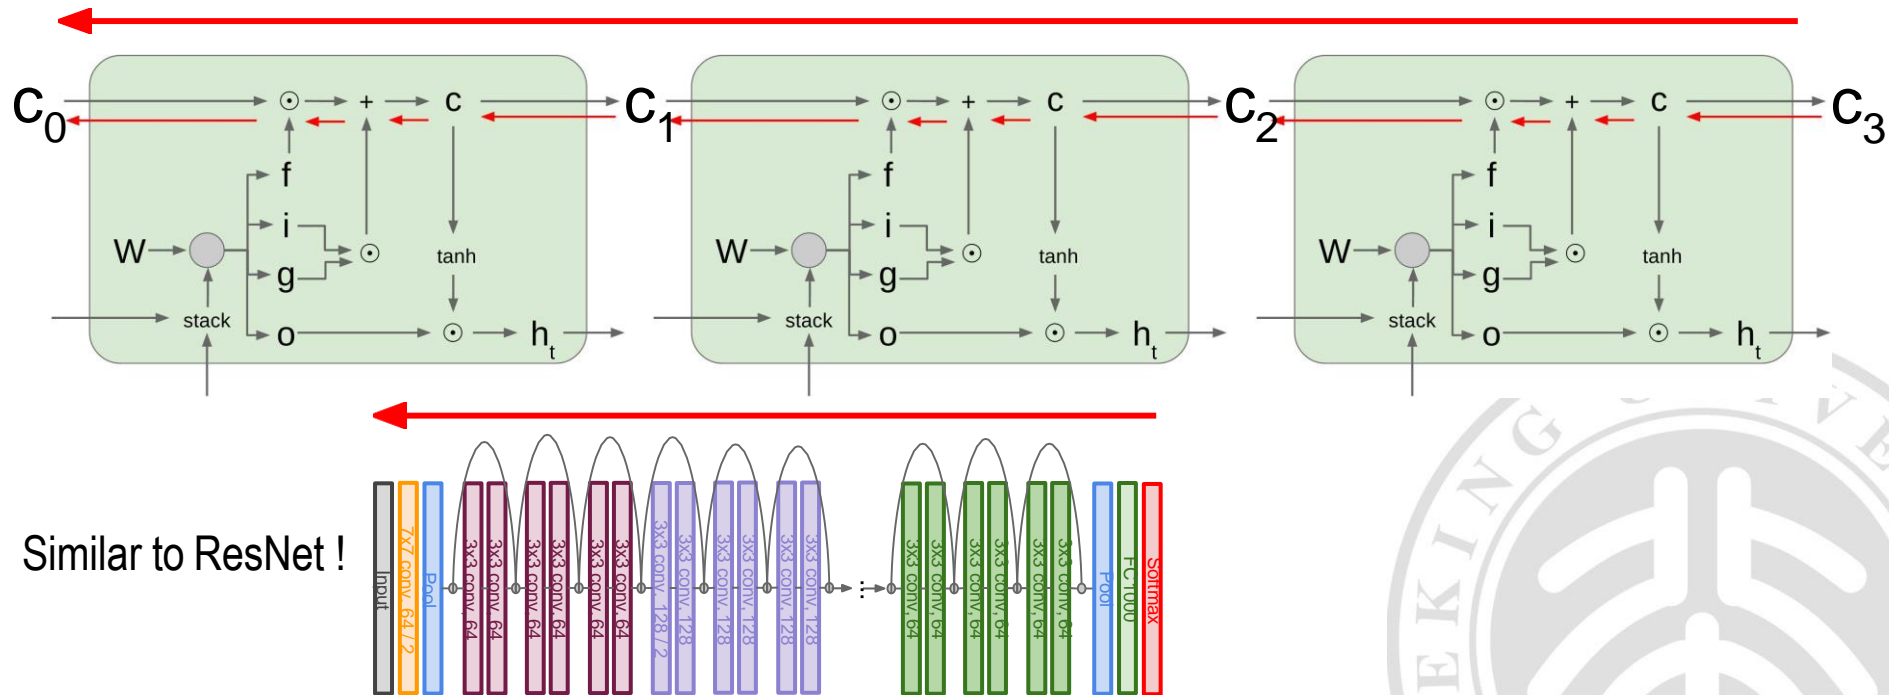
LSTM

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$

Uninterrupted gradient flow!



Similar to ResNet !

- CNN Architectures
 - AlexNet
 - VGG
 - GoogLeNet
 - ResNet
- Recurrent Neural Network
 - Vanilla RNN
 - Backpropagation through time
 - Long Short-Term Memory
- **Beyond CNN and RNN**
 - **Unsupervised Learning**
 - **Generative Adversarial Network**



Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a *function* to map $x \rightarrow y$

Examples: Classification, regression, object detection, semantic segmentation, image captioning, etc.



→ Cat

Classification



Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a *function* to map $x \rightarrow y$

Examples: Classification, regression, object detection, semantic segmentation, image captioning, etc.



DOG, DOG, CAT

Object Detection


Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a *function* to map $x \rightarrow y$

Examples: Classification, regression, object detection, semantic segmentation, image captioning, etc.



GRASS, CAT,
TREE, SKY

Semantic Segmentation

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a *function* to map $x \rightarrow y$

Examples: Classification, regression, object detection, semantic segmentation, image captioning, etc.



A cat sitting on a suitcase on the floor

Image Captioning


Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.



K-means clustering

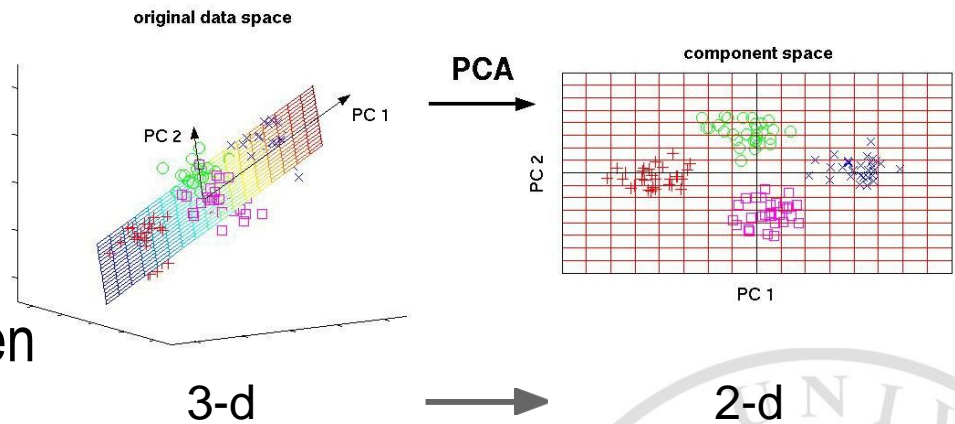
Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.



Principal Component Analysis
(Dimensionality reduction)

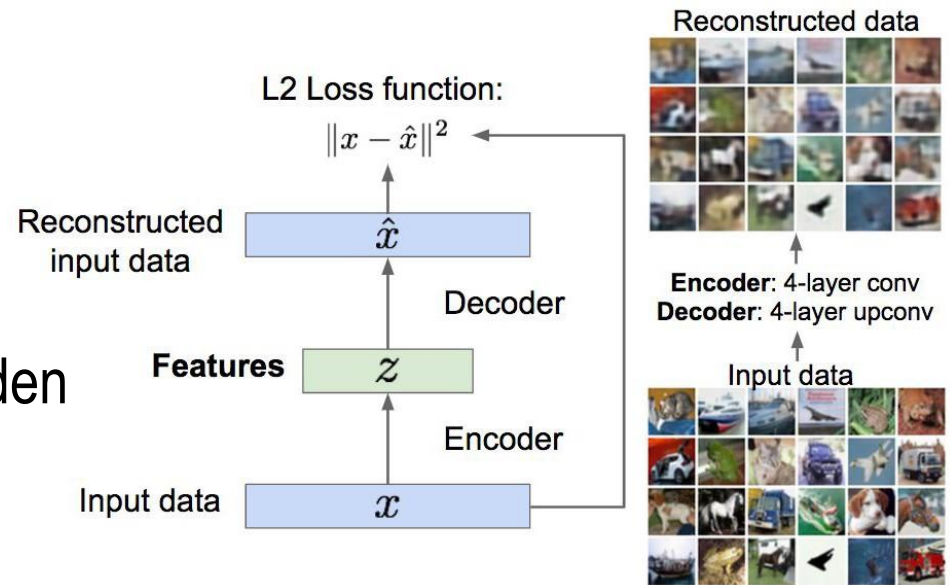
Unsupervised Learning

Data: x

Just data, no labels!


Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.




Autoencoders
(Feature learning)

- Given training data, generate new samples from same distribution
- **Several flavors:**
 - Explicit density estimation: explicitly define and solve for $p_{\text{model}}(x)$
 - Implicit density estimation: learn model that can sample from $p_{\text{model}}(x)$ w/o explicitly defining it



Training data $\sim p_{\text{data}}(x)$




Generated samples $\sim p_{\text{model}}(x)$


Want to learn $p_{\text{model}}(x)$ similar to $p_{\text{data}}(x)$

Why Generative Models

- Realistic samples for artwork, super-resolution, colorization, etc.




- Generative models of time-series data can be used for simulation and planning (reinforcement learning applications!)
- Training generative models can also enable inference of latent representations that can be useful as general features



- Problem: Want to sample from complex, high-dimensional training distribution. No direct way to do this!
- Solution: Sample from a simple distribution, e.g. random noise. Learn transformation to training distribution.
- Q: What can we use to represent this complex transformation?
- A: A neural network!

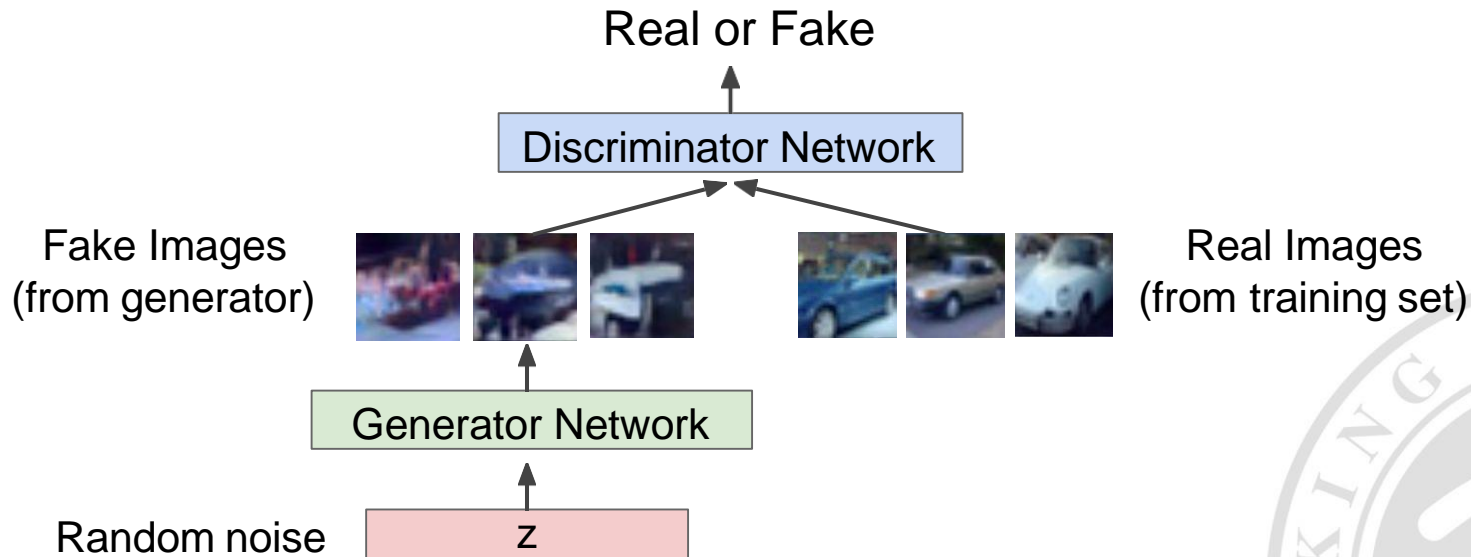
Output: Sample from training distribution

Input: Random noise



Generator network: try to fool the discriminator by generating real-looking images

Discriminator network: try to distinguish between real and fake images



Training GANs: Two-player game

Generator network: try to fool the discriminator by generating real-looking images

Discriminator network: try to distinguish between real and fake images

Train jointly in **minimax game**


Minimax objective function:

Discriminator outputs likelihood in (0,1) of real image


$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\substack{\text{Discriminator output} \\ \text{for real data } x}} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\substack{\text{Discriminator output for} \\ \text{generated fake data } G(z)}}) \right]$$

- Discriminator (θ_d) wants to **maximize objective** such that $D(x)$ is close to 1 (real) and $D(G(z))$ is close to 0 (fake)
- Generator (θ_g) wants to **minimize objective** such that $D(G(z))$ is close to 1 (discriminator is fooled into thinking generated $G(z)$ is real)

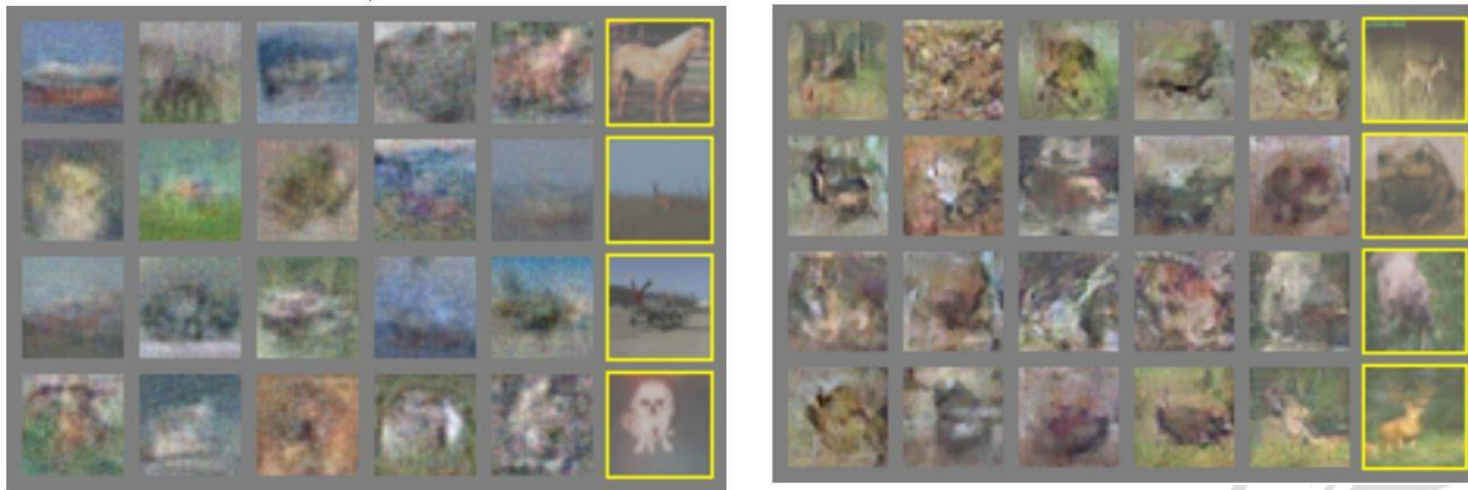
Generated samples



Nearest neighbor from training set




Generated samples (CIFAR-10)



Nearest neighbor from training set

Generator is an upsampling network with fractionally-strided convolutions
Discriminator is a convolutional network

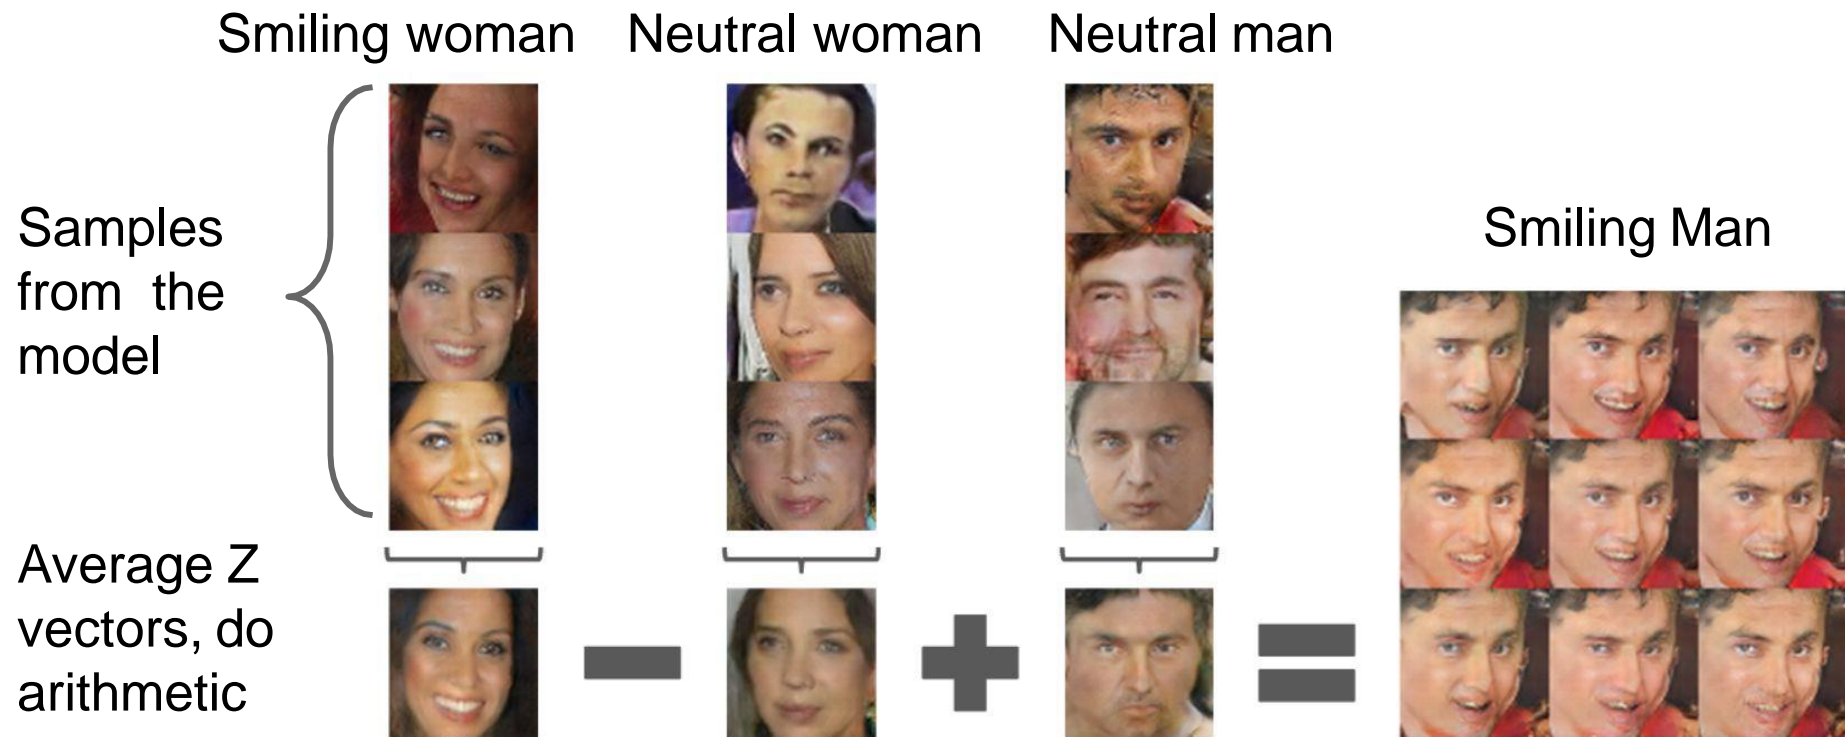


Generator

Samples from the model look much better!



Interpretable Vector Math



- Don't work with an explicit density function
- Take game-theoretic approach: learn to generate from training distribution through 2-player game
- Pros:
 - Beautiful, state-of-the-art samples!
- Cons:
 - Trickier / more unstable to train
- Active areas of research:
 - Better loss functions, more stable training (Wasserstein GAN, LSGAN, many others)
 - Conditional GANs, GANs for all kinds of applications

